

# Parallel sorting algorithms

## Theory and implementation

Xavier JUVIGNY, AKOU, DAAA, ONERA

[xavier.juvigny@onera.fr](mailto:xavier.juvigny@onera.fr)

Course Parallel Programming

- February 3rd 2025 -

<sup>1</sup> ONERA, <sup>2</sup> DAAA

# Table of contents

---

- 1 Theory of parallel sorting algorithms
- 2 Parallel sort algorithms
- 3 Quicksort algorithm
- 4 Bitonic sorting algorithm
- 5 Bucket-sort algorithms

# Overview

---

1 Theory of parallel sorting algorithms

2 Parallel sort algorithms

3 Quicksort algorithm

4 Bitonic sorting algorithm

5 Bucket-sort algorithms

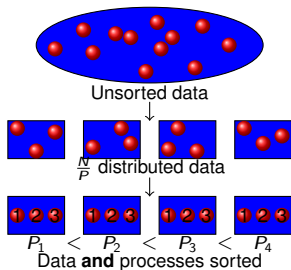
# Complexity of sorting algorithms

## Basic operations

- **Compare algorithm** : Comparison algorithm complexity is supposed  $\mathcal{O}(1)$ . But in distributed parallel context, one must consider the distribution of the initial data to account for the cost of data exchange between processes !
- **Exchange algorithm** : Exchange algorithm complexity is supposed  $\mathcal{O}(1)$ . But same consideration to do as **compare algorithm** ;
- Sequential “compare-and-exchange” algorithm :

```
if (a>b) { // Comparison
  // Exchange
  tmp = a;
  a = b;
  b = tmp; }
```

# Potential speed-up



- Best sequential sorting algorithms ( for arbitrary sequences of numbers ) have average time complexity  $O(n \log n)$
- hence, the best speedup one can expect from using  $n$  processors is  $\frac{O(n \log n)}{n} = O(\log n)$
- there are such parallel algorithms, but the hidden constant is very large ( F. Thomson Leighton : Introduction to parallel algorithms and architectures (1991) )
- In general, a practical useful  $O(\log n)$  algorithm may be difficult to find.

**Beware**, it may be a bad idea to take  $n$  processes to sort  $n$  data (granularity).

# Parallelization of a naive algorithm

## Naive algorithm

- Count the number of numbers that are smaller than a number  $a$  in the list
- this gives the position of  $a$  in the sorted list
- this procedure has to be repeated for all elements of the list; hence the time complexity is  $n(n-1) = O(n^2)$  ( not so good sequential algorithm )

## Implementation

```
for ( i = 0; i < n; i++ ) { // For each value
    x = 0;
    for ( j = 0; j < n; j++ ) // Computing the new pos.
        if ( a[i] > a[j] ) x++;
    b[x] = a[i];
}
```

Works well if there are no repetitions of the numbers in the list (in case of repetitions the code must be changed slightly).

# Rank sort : Parallel code

## Embarrassingly “ideal” algorithm

Parallel code, using  $n$  processes ( for  $n$  values to sort )

```
x = 0;  
for ( j = 0; j < n; j++ )  
    if ( a[rank] > a[j] ) x++;  
b[x] = a[rank];
```

## Complexity

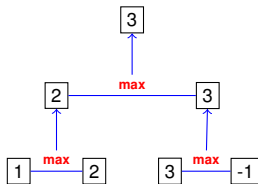
- $n$  processors work in parallel to find the ranks of all numbers of the list ;
- Parallel time complexity is  $O(n)$ , better than any sequential sorting algorithm !
- Usable for GPGPU units.

# More parallelization...

Parallel code using  $n^2$  processes ( for  $n$  values to sort )

## Parallel algorithm

- In the case  $n^2$  processes may be used, the comparison of each  $a[0], \dots, a[n-1]$  with  $a[i]$  may be done in parallel as well
- Incrementing the counter is still sequential, hence the overall computation requires  $1 + n$  steps ;
- If a tree structure is used to increment the counter, then the overall computation time is  $O(\log_2 n)$



( but, as one expects, processor efficiency is very low )

These are just theoretical results : it is not efficient to use  $n$  or  $n^2$  processors to sort  $n$  numbers.



# Data partitioning

## Context

- Usually the number  $n$  of values is much larger than the number  $p$  of processes ;
- In such cases, each process will handle a part of the data (a sublist of the data)

## Distributed sorted container

- local container is sorted ;
- if  $p_i < p_j$  then  $\forall a_i \in p_i, \forall a_j \in p_j, a_i \leq a_j$

## Global scheme of parallel sort algorithm

For a process :

- Sort his local data ;
- Run a merge sort algorithm to concatenate its list with that received from another process ;
- Keep the bottom half (or the top half) of the sorted list.

# Parallel compare and exchange operations

## Asymmetric algorithm

- Process  $p_i$  sends local value  $A$  to process  $p_j$  ;
- Process  $p_j$  compares value  $A$  with some local values  $B_j$  ;
- Send the  $B_j$  which are larger (or lesser) than  $A$ . If no  $B_j$  is larger (or lesser) than  $A$ , send back  $A$  ;

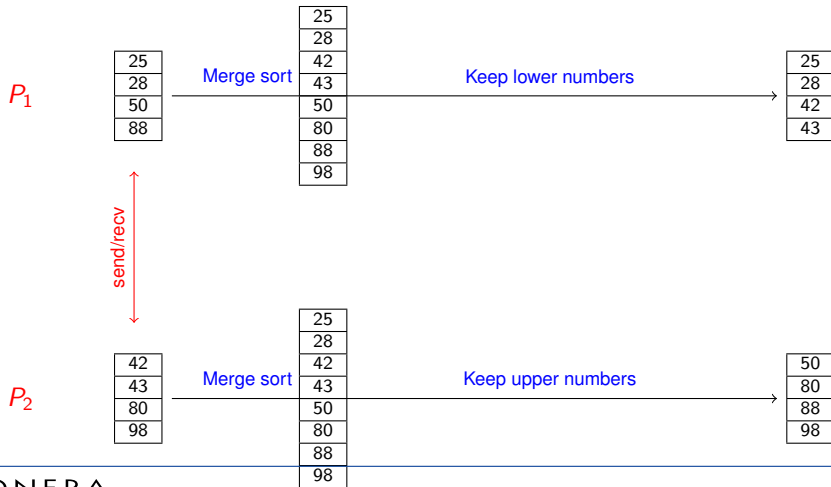
## Symmetric algorithm

- Processes  $p_i$  and  $p_j$  send some value to the other ;
- Each process compares his value with the received value ;
- Each process keeps his value or the received value relative to the comparison result ;

## Remarks

- Data exchanges between processes is very expensive, so find some algorithms which minimize data exchanges ;
- In general, the receive operation doesn't know the number of values to receive  $\Rightarrow$  one must probe the received message to get the number of data to receive, allocate the relative buffer and receive the data !

# Scheme of a general algorithm for parallel sort algorithm



# Overview

---

- 1 Theory of parallel sorting algorithms
- 2 Parallel sort algorithms

- 3 Quicksort algorithm
- 4 Bitonic sorting algorithm
- 5 Bucket-sort algorithms

# Sequential bubble sort algorithm

## Bubble sort algorithm

- Simplest, but not an efficient sequential sorting algorithm ;
- Compare/exchange complexity :

$$\sum_{i=1}^{n-1} i = \frac{n(n-1)}{2} = O(n^2)$$



Figure – Analogic bubble sort

## Sequential code

```
for (int i=n-1; i>0; --i)
  for (int j=0; j<i; ++j) {
    k = j+1;
    if (a[j]>a[k]) std::swap(a[j],a[k]);
  }
```

# Odd-Even sort algorithm

- Parallelized bubble sort
- Based on idea that the bodies of the main loop may be overlapped

## “scalar” Algorithm : Iteration between even and odd phase

- Even phase



```
if (rank%2==0) {  
    recv(&temp, (rank+1)%nbp);  
    send(&value, (rank+1)%nbp);  
    if (temp < A) A = temp; }  
}
```

```
if (rank%2==1) {  
    send(&value, rank-1);  
    recv(&temp, rank-1);  
    if (temp > A) A = temp; }  
}
```

- Odd phase



```
if (rank%2==0) {  
    recv(&temp, (rank+nbp-1)%nbp);  
    send(&value, (rank+nbp-1)%nbp);  
    if (temp > A) A = temp; }  
}
```

```
if (rank%2==1) {  
    send(&value, (rank+1)%nbp);  
    recv(&temp, (rank+1)%nbp);  
    if (temp < A) A = temp; }  
}
```

# Example of even-odd parallel bubble sort

Example : Sorting 8 numbers on 8 processes

Step	$P_0$		$P_1$		$P_2$		$P_3$		$P_4$		$P_5$		$P_6$		$P_7$
0	4	↔	2		7	↔	8		5	↔	1		3	↔	6
1	2		4	↔	7		8	↔	1		5	↔	3		6
2	2	↔	4		7	↔	1		8	↔	3		5	↔	6
3	2		4	↔	1		7	↔	3		8	↔	5		6
4	2	↔	1		4	↔	3		7	↔	5		8	↔	6
5	1		2	↔	3		4	↔	5		7	↔	6		8
6	1	↔	2		3	↔	4		5	↔	6		7	↔	8
7	1		2	↔	3		4	↔	5		6	↔	7		8

# Odd-even parallel algorithm per block

## Per block algorithm

- Replace a value per process with a sorted set of values per process
- Use sort-fusion algorithm to exchange values
- Data comparison complexity :  
$$\frac{N}{nbp} \log_2 \left( \frac{N}{nbp} \right) + (nbp - 1) \cdot \frac{2N}{nbp}$$
- Data communication complexity :  
$$(nbp - 1) \cdot \frac{2N}{nbp}$$

## Implementation

```
sort(values); // Quick sort of local values
for (iter=0; iter<nbp-1; ++iter) { // Odd-even algorithm
    if (iter is odd) {
        if (rank is even and rank > 0) {
            recv(buffer, rank-1); send(values, rank-1);
            values = fusionSort(buffer, values, keepMax);
        } else if (rank is odd and rank < nbp-1) {
            send(values, rank+1); recv(buffer, rank+1);
            values = fusionSort(buffer, values, keepMin);
        }
    } else if (iter is even) {
        if (rank is even and rank < nbp-1) {
            recv(buffer, rank+1); send(values, rank+1);
            values = fusionSort(buffer, values, keepMin);
        } else if (rank is odd) {
            send(values, rank-1); recv(buffer, rank-1);
            values = fusionSort(buffer, values, keepMax);
        }
    }
}
```



# Shear sort algorithm

## Two dimensional sorting

### Basic idea

- Look at the array as a two-dimensional array (one row per process)
- The goal is to sort this 2D array in snakelike style : even rows increasing, odd rows decreasing ;
- Two phases : In **even phase**, sort per row, in **odd phase**, sort per column increasing from top to bottom ;
- After  $\log_2(N) + 1$  phases, the array is snakelike-style sorted.

### Example

4	14	8	2
10	3	13	16
7	15	1	5
12	6	11	9

Original number

### Remarks

- Embarrassingly parallel algorithm for shared memory !
- But not well adapted for distributed parallel architecture as is ;
- How to change this algorithm for distributed parallel architecture ?

# Shear sort algorithm

## Two dimensional sorting

### Basic idea

- Look at the array as a two-dimensional array (one row per process)
- The goal is to sort this 2D array in snakelike style : even rows increasing, odd rows decreasing ;
- Two phases : In **even phase**, sort per row, in **odd phase**, sort per column increasing from top to bottom ;
- After  $\log_2(N) + 1$  phases, the array is snakelike-style sorted.

### Example

2	4	8	14
16	13	10	3
1	5	7	15
12	11	9	6

Phase 1 – row sort

### Remarks

- Embarrassingly parallel algorithm for shared memory !
- But not well adapted for distributed parallel architecture as is ;
- How to change this algorithm for distributed parallel architecture ?

# Shear sort algorithm

## Two dimensional sorting

### Basic idea

- Look at the array as a two-dimensional array (one row per process)
- The goal is to sort this 2D array in snakelike style : even rows increasing, odd rows decreasing ;
- Two phases : In **even phase**, sort per row, in **odd phase**, sort per column increasing from top to bottom ;
- After  $\log_2(N) + 1$  phases, the array is snakelike-style sorted.

### Example

1	4	7	3
2	5	8	6
12	11	9	14
16	13	10	15

Phase 2 – col sort

### Remarks

- Embarrassingly parallel algorithm for shared memory !
- But not well adapted for distributed parallel architecture as is ;
- How to change this algorithm for distributed parallel architecture ?

# Shear sort algorithm

## Two dimensional sorting

### Basic idea

- Look at the array as a two-dimensional array (one row per process)
- The goal is to sort this 2D array in snakelike style : even rows increasing, odd rows decreasing ;
- Two phases : In **even phase**, sort per row, in **odd phase**, sort per column increasing from top to bottom ;
- After  $\log_2(N) + 1$  phases, the array is snakelike-style sorted.

### Example

1	3	4	7
8	6	5	2
9	11	12	14
16	15	13	10

Phase 3 : row sort

### Remarks

- Embarrassingly parallel algorithm for shared memory !
- But not well adapted for distributed parallel architecture as is ;
- How to change this algorithm for distributed parallel architecture ?

# Shear sort algorithm

## Two dimensional sorting

### Basic idea

- Look at the array as a two-dimensional array (one row per process)
- The goal is to sort this 2D array in snakelike style : even rows increasing, odd rows decreasing ;
- Two phases : In **even phase**, sort per row, in **odd phase**, sort per column increasing from top to bottom ;
- After  $\log_2(N) + 1$  phases, the array is snakelike-style sorted.

### Example

1	3	4	2
8	6	5	7
9	11	12	10
16	15	13	14

Phase 4 : col sort

### Remarks

- Embarrassingly parallel algorithm for shared memory !
- But not well adapted for distributed parallel architecture as is ;
- How to change this algorithm for distributed parallel architecture ?

# Shear sort algorithm

## Two dimensional sorting

### Basic idea

- Look at the array as a two-dimensional array (one row per process)
- The goal is to sort this 2D array in snakelike style : even rows increasing, odd rows decreasing ;
- Two phases : In **even phase**, sort per row, in **odd phase**, sort per column increasing from top to bottom ;
- After  $\log_2(N) + 1$  phases, the array is snakelike-style sorted.

### Example

1	2	3	4
8	7	6	5
9	10	11	12
16	15	14	13

Final 5 : row sort

### Remarks

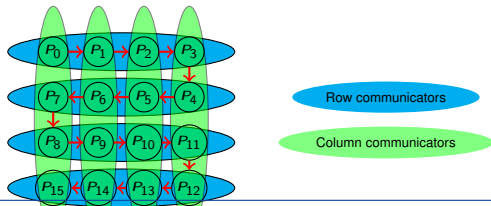
- Embarrassingly parallel algorithm for shared memory !
- But not well adapted for distributed parallel architecture as is ;
- How to change this algorithm for distributed parallel architecture ?

# Shear sort algorithm for parallel distributed memory architecture

## Implementation ideas

- Same principal as odd-even algorithm : replace a value with some sets of values  $S_i$  (one set per process) ;
- Define a relation order :  $S_i < S_j$  iff  $\max(S_i) < \min(S_j)$  (In set, values ordered as increasing order)
- Use odd-even algorithm to parallelize the phase of sorting per row or column ;
- Grouping processes in new communicators per rows and per columns ;
- Play with rank numbering to alternate between increasing order and decreasing order for rows ;

## Processes repartition



- Use `MPI_Comm_split(comm,color,key,&newcomm)` to define row and columns communicators ;
- Processes calling this function with same color are inside the same new communicator ;
- key is a value to number the processes inside the new communicator.

# Overview

---

- 1 Theory of parallel sorting algorithms
- 2 Parallel sort algorithms

- 3 Quicksort algorithm
- 4 Bitonic sorting algorithm
- 5 Bucket-sort algorithms



# Sequential quick-sort algorithm

## Reminder

- Optimal sequential sorting algorithm ( $\mathcal{O}(n \log_2(n))$ ) based on divide-and-conquer algorithm class ;
- Select a number  $r$  called **pivot** and split the list into two sublists : one with all elements at most equal to  $r$ , the other holding all elements greater than  $r$  ;
- This procedure is recursive, applied until one element lists are obtained (which are sorted)
- Example :

4   2   7   8   5   1   3   6

## Sequential code

```
void quicksort( T* list, T* start, T* end ) {  
    auto pivot = choosePivot(start,end);  
    if (start < end) {  
        split(list, start, end, pivot);  
        quicksort(list, start, pivot-1);  
        quicksort(list,pivot+1, end );  
    }  
}
```

# Sequential quick-sort algorithm

## Reminder

- Optimal sequential sorting algorithm ( $\mathcal{O}(n \log_2(n))$ ) based on divide-and-conquer algorithm class ;
- Select a number  $r$  called **pivot** and split the list into two sublists : one with all elements at most equal to  $r$ , the other holding all elements greater than  $r$  ;
- This procedure is recursive, applied until one element lists are obtained (which are sorted)
- Example :

4   2   7   8   5   1   3   6

## Sequential code

```
void quicksort( T* list, T* start, T* end ) {  
    auto pivot = choosePivot(start,end);  
    if (start < end) {  
        split(list, start, end, pivot);  
        quicksort(list, start, pivot-1);  
        quicksort(list,pivot+1, end );  
    }  
}
```

# Sequential quick-sort algorithm

## Reminder

- Optimal sequential sorting algorithm ( $\mathcal{O}(n \log_2(n))$ ) based on divide-and-conquer algorithm class ;
- Select a number  $r$  called **pivot** and split the list into two sublists : one with all elements at most equal to  $r$ , the other holding all elements greater than  $r$  ;
- This procedure is recursive, applied until one element lists are obtained (which are sorted)
- Example :

4   2   5   1   3   6   7   8

## Sequential code

```
void quicksort( T* list, T* start, T* end ) {  
    auto pivot = choosePivot(start,end);  
    if (start < end) {  
        split(list, start, end, pivot);  
        quicksort(list, start, pivot-1);  
        quicksort(list,pivot+1, end );  
    }  
}
```

# Sequential quick-sort algorithm

## Reminder

- Optimal sequential sorting algorithm ( $\mathcal{O}(n \log_2(n))$ ) based on divide-and-conquer algorithm class ;
- Select a number  $r$  called **pivot** and split the list into two sublists : one with all elements at most equal to  $r$ , the other holding all elements greater than  $r$  ;
- This procedure is recursive, applied until one element lists are obtained (which are sorted)
- Example :

4   2   5   1   3   6   7   8

## Sequential code

```
void quicksort( T* list, T* start, T* end ) {  
    auto pivot = choosePivot(start,end);  
    if (start < end) {  
        split(list, start, end, pivot);  
        quicksort(list, start, pivot-1);  
        quicksort(list,pivot+1, end    );  
    }  
}
```

# Sequential quick-sort algorithm

## Reminder

- Optimal sequential sorting algorithm ( $\mathcal{O}(n \log_2(n))$ ) based on divide-and-conquer algorithm class ;
- Select a number  $r$  called **pivot** and split the list into two sublists : one with all elements at most equal to  $r$ , the other holding all elements greater than  $r$  ;
- This procedure is recursive, applied until one element lists are obtained (which are sorted)
- Example :

2   1   3   4   5   6   7   8

## Sequential code

```
void quicksort( T* list, T* start, T* end ) {  
    auto pivot = choosePivot(start,end);  
    if (start < end) {  
        split(list, start, end, pivot);  
        quicksort(list, start, pivot-1);  
        quicksort(list,pivot+1, end    );  
    }  
}
```

# Sequential quick-sort algorithm

## Reminder

- Optimal sequential sorting algorithm ( $\mathcal{O}(n \log_2(n))$ ) based on divide-and-conquer algorithm class ;
- Select a number  $r$  called **pivot** and split the list into two sublists : one with all elements at most equal to  $r$ , the other holding all elements greater than  $r$  ;
- This procedure is recursive, applied until one element lists are obtained (which are sorted)
- Example :

2   1   3   4   5   6   7   8

## Sequential code

```
void quicksort( T* list, T* start, T* end ) {  
    auto pivot = choosePivot(start,end);  
    if (start < end) {  
        split(list, start, end, pivot);  
        quicksort(list, start, pivot-1);  
        quicksort(list,pivot+1, end    );  
    }  
}
```

# Sequential quick-sort algorithm

## Reminder

- Optimal sequential sorting algorithm ( $\mathcal{O}(n \log_2(n))$ ) based on divide-and-conquer algorithm class ;
- Select a number  $r$  called **pivot** and split the list into two sublists : one with all elements at most equal to  $r$ , the other holding all elements greater than  $r$  ;
- This procedure is recursive, applied until one element lists are obtained (which are sorted)
- Example :

1 2 3 4 5 6 7 8

## Sequential code

```
void quicksort( T* list, T* start, T* end ) {  
    auto pivot = choosePivot(start,end);  
    if (start < end) {  
        split(list, start, end, pivot);  
        quicksort(list, start, pivot-1);  
        quicksort(list,pivot+1, end );  
    }  
}
```

# Naive parallelization of quicksort algorithm

## Ideas

- The class of quick-sort algorithm suggests to apply a divide-and-conquer parallelization method ;
- The main problem of this approach is that the tree distribution ( induced by the lengths of the sublists) heavily depends on pivot selection ; in the worst case, the tree may consists of a single path (as a list) ;
- **Analysis** : provided an equal distribution of values within sublists is ensured, one gets :
  - Comparisons :  $n + \frac{n}{2} + \frac{n}{4} + \dots \approx 2n$
  - Communications :  $t_s + \frac{n}{2}t_d + t_s + \frac{n}{4}t_d + \dots \approx \log_2(n)t_s + nt_d$  where  $t_s$  is time to start a communication and  $t_d$  the time to transfer one element to another process.
- But only last iteration uses full parallelization !



# Hyperquick sort algorithm

---

Binary numbering of Hypercube

0

Figure – Dimension 0

# Hyperquick sort algorithm

---

Binary numbering of Hypercube



Figure – Dimension 1

# Hyperquick sort algorithm

Binary numbering of Hypercube

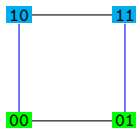


Figure – Dimension 2

# Hyperquick sort algorithm

Binary numbering of Hypercube

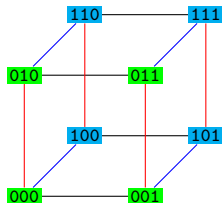


Figure – Dimension 3

# Hyperquick sort algorithm

Binary numbering of Hypercube




Figure – Dimension 4

# Hyper quick sort algorithm (2)

## Numbering vertices of hyper-cube

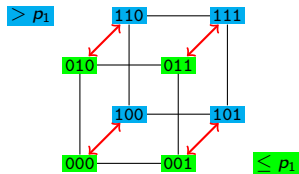
- The binary numbers of two linked nodes have a difference of one bit ;
- The distance between two nodes in a hypercube ( minimal number of nodes to access to go from first node to second node ) is the number of bit which differ in their binary number ;
- It's the **Gray code** numbering.

## Ideas of the hyper quick sort algorithm


- Initially, Data are distributed across all processes ;
- Each process sorts its local data ;
- Loop on dimension of the hypercube and for dimension  $d$ , consider pair of processes  $(p; p + 2^d)$
- Process  $p$  chooses his median value as pivot and sends value lesser than pivot to process  $p + 2^d$  ;
- Process  $p + 2^d$  receives pivot and data from process  $p$  and keeps value greater than pivot ;
- Each process keeps sorted value, using fusion sort algorithm to keep sorting.

# Illustration of hyper quick sort

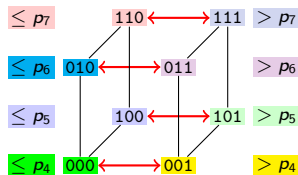
Phase 1 :



Phase 2 :



Phase 3 :



# Hyperquicksort : complexity analysis

- Suppose we run algorithm on  $nbp = 2^d$  processes (hypercube dimension  $d$ );
- Each process holds initially  $N_I = \frac{N}{nbp}$  values;
- **Initial sorting** :  $N_I \cdot \log_2(N_I)$  comparisons;
- **Pivot selection** :  $\mathcal{O}(1)$  (one takes middle list element) for each dimension;
- **Pivot broadcasting** :
  - Broadcast one pivot in a  $k$  hypercube :  $k(t_s + t_d)$ ;
  - For all iterations :  $(d + \dots + 1)(t_s + t_d) = \frac{d(d-1)}{2}(t_s + t_d)$  comparisons;
- **Split list from pivot value** : For sorted list of size  $x$ ,  $\log_2(x)$  comparisons;
- **Exchange part of list** : To exchange  $\frac{x}{2}$  data :  $2(t_s + \frac{x}{2}t_d)$
- **Fusion merge sort** :  $\frac{x}{2}$  comparisons

## Total (for ideal balance)

- $N_I \cdot \log_2(N_I) + \left(\log_2(N_I) + \frac{N_I}{2}\right) \cdot d$  comparisons
- $\left(\frac{d(d-1)}{2} + 2d\right) t_s + \left(\frac{d(d-1)}{2} + d \cdot N_I\right) \cdot t_d$  for message;



# Overview

---

- 1 Theory of parallel sorting algorithms
- 2 Parallel sort algorithms

- 3 Quicksort algorithm
- 4 Bitonic sorting algorithm
- 5 Bucket-sort algorithms

# Bitonic sequences

## Definition of a bitonic sequence

- A sequence of values  $\{a_i\}_{i \in [1;N]}$  than can be split in two subsequences (of consecutive numbers), one increasing and one decreasing ; e.g :

$$\exists i \in [1; N] \text{ verifying } \begin{cases} a_1 \leq a_2 \leq \dots \leq a_i \\ a_i \geq a_{i+1} \geq \dots \geq a_N \end{cases}$$

Example : 3, 5, 8, 19, 17, 14, 12, 11

- Or** a sequence which may be brought to this form by a circular shifting of the elements of the sequence

Example : 12, 11, 3, 5, 8, 19, 17, 14

## Remarks

- The first subsequence can be increasing **or** decreasing.
- So a bitonic sequence (without considering circular shifting) can be increasing-decreasing or decreasing-increasing.
- A monotone bitonic sequence is a sorted sequence (increasing or decreasing) ;
- All sequences with three or less elements are bitonic.

# Splitting a bitonic sequence

## Bitonic split

Let  $\{a_i\}_{i \in [1;N]}$  be a bitonic sequence. So the subsequences

$$\begin{cases} \{b_i\}_{i \in [1; \frac{N}{2}]} \\ \{c_i\}_{i \in [1; \frac{N}{2}]} \end{cases} = \begin{cases} \min(a_1, a_{1+\frac{N}{2}}), \min(a_2, a_{2+\frac{N}{2}}), \dots, \min(a_{\frac{N}{2}-1}, a_N) \\ \max(a_1, a_{1+\frac{N}{2}}), \max(a_2, a_{2+\frac{N}{2}}), \dots, \max(a_{\frac{N}{2}-1}, a_N) \end{cases}$$

- are bitonic sequences ;
- $\forall i \in [1; \frac{N}{2}] ; b_i \leq c_i$

## Example

1 5 8 7 **6 4 3 2**  $\xRightarrow{\text{split}}$  1 4 3 2 **6 5 8 7**

# Sorting a bitonic sequence (SBS algorithm)

## Algorithm

Apply split procedure on bitonic sequence and repeat this splitting procedure on subsequences until having one element per subsequence to obtain sorted sequence.

## Example

1

1 5 8 7 6 4 3 2

2

1 4 3 2 6 5 8 7

3

1 2 3 4 6 5 8 7

4

1 2 3 4 5 6 7 8

# Sorting a bitonic sequence (SBS algorithm)

## Algorithm

Apply split procedure on bitonic sequence and repeat this splitting procedure on subsequences until having one element per subsequence to obtain sorted sequence.

## Example

1

1 5 8 7 6 4 3 2

2

1 4 3 2 6 5 8 7

3

1 2 3 4 6 5 8 7

4

1 2 3 4 5 6 7 8

# Sorting a bitonic sequence (SBS algorithm)

## Algorithm

Apply split procedure on bitonic sequence and repeat this splitting procedure on subsequences until having one element per subsequence to obtain sorted sequence.

## Example

1

1 5 8 7 6 4 3 2

2

1 4 3 2 6 5 8 7

3

1 2 3 4 6 5 8 7

4

1 2 3 4 5 6 7 8

# Sorting a bitonic sequence (SBS algorithm)

## Algorithm

Apply split procedure on bitonic sequence and repeat this splitting procedure on subsequences until having one element per subsequence to obtain sorted sequence.

## Example

1

1 5 8 7 6 4 3 2

2

1 4 3 2 6 5 8 7

3

1 2 3 4 6 5 8 7

4

1 2 3 4 5 6 7 8

# Building a bitonic sequence

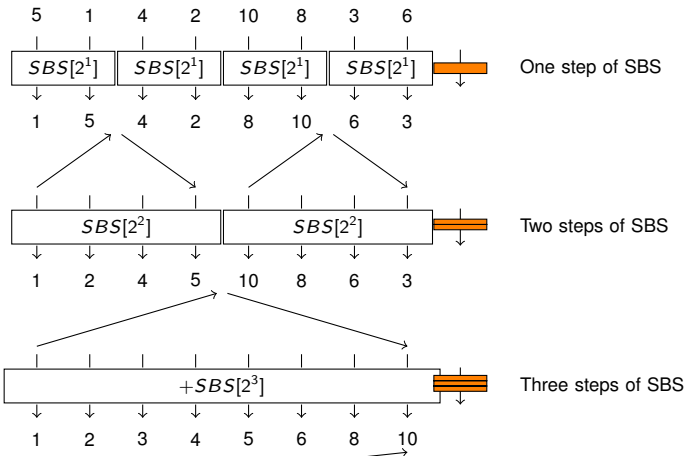
---

## Procedure

- 1 Split the sequence in two-elements subsequences (which are bitonics !);
- 2 Sort subsequences alternating increasing and decreasing sorting (using SBS algorithm);
- 3 Concatenate two adjacent lists to get a longer bitonic sequence;
- 4 Repeat from step 2 until the full list becomes a bitonic sequence.



# Example on 8 elements



# Bitonic sort analysis

## Complexity of bitonic sort :

- With  $n = 2^k$ , there are  $k$  phases, each involving  $1, 2, \dots, k$  steps, respectively ;
- The total number of steps is

$$\sum_{i=1}^k i = \frac{k(k+1)}{2} = O(k^2) = O(\log_2^2 n) \quad (1)$$

- Total complexity is  $O(n \log_2^2 n)$

# Adapt bitonic sort to distributed parallel architecture

---

## Main ideas

- First sort local data with the fastest sort algorithm in increasing values if rank is even and decreasing values if rank is odd ;
- Pairing inside a subcommunicator processes to define a bitonic sequence ;
- And apply bitonic sort algorithm, grouping processes per four, height and so on. . .
- Sub-communicators built here are very similar to the sub-communicators build with hyperquick sort algorithms !

# Overview

---

- 1 Theory of parallel sorting algorithms
- 2 Parallel sort algorithms

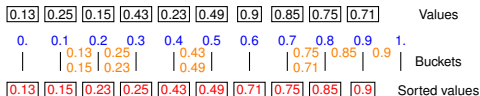
- 3 Quicksort algorithm
- 4 Bitonic sorting algorithm
- 5 Bucket-sort algorithms

# Bucket Sort algorithm

## Algorithm

Distribute elements of a list into a fixed number of buckets, and sort the elements inside each buckets.

- Setup an array of initially empty "buckets"
- **Scatter** : Put each element of the list in the right bucket
- Sort the elements inside each bucket
- **Gather** : Visit each bucket in order and put all elements back in original list



- Complexity at best :  $N + \frac{N}{k} \log_2(\frac{N}{k})$  comparisons ( $k$  = numbers of buckets)
- **Difficulty** : How to choose the intervals for the buckets ?

# Parallel bucket sort

## Main ideas for the parallel algorithm

- Each process is seen as one bucket ;
- Compute intervals for buckets :
  - Sort local data
  - Take  $nbp + 1$  values at regular intervals
  - Gather values in bucket array
  - Extract  $nbp + 1$  values to find intervals for buckets

## Analysis

- Local sort :  $\frac{N}{nbp} \log_2(\frac{N}{nbp})$  comparisons ;
- Gather bucket array :  $nbp(t_s + (nbp + 1)t_d)$
- Sort bucket array :  $\frac{nbp}{\log_2} (nbp)$  comparisons ;
- Distribute data in buckets :  $(nbp - 1) \left( t_s + \frac{N}{nbp^2} t_d \right)$
- Sort local data :  $2(nbp - 1) \frac{N}{nbp}$  comparisons (fusion sort)